

Use Delta Tables in Apache Spark

[Create a workspace](#)

[Create a lakehouse and upload files](#)

[Explore data in a DataFrame](#)

[Create Delta tables](#)

[Use SQL to create a Delta table](#)

[Explore table versioning](#)

[Analyze Delta table data with SQL queries](#)

[Use Delta tables for streaming data](#)

[Clean up resources](#)

Tables in a Microsoft Fabric Lakehouse are based on the open-source Delta Lake format. Delta Lake adds support for relational semantics for both batch and streaming data. In this exercise you will create Delta tables and explore the data using SQL queries.

This exercise should take approximately **45** minutes to complete

! [!Note] You need access to a [Microsoft Fabric tenant](#) to complete this exercise.

Create a workspace


Before working with data in Fabric, create a workspace in a tenant with the Fabric capacity enabled.

1. Navigate to the [Microsoft Fabric home page](#) at

`https://app.fabric.microsoft.com/home?experience=fabric-developer` in a browser and sign in with your Fabric credentials.

2. In the menu bar on the left, select **Workspaces** (the icon looks similar to ).
3. Create a new workspace with a name of your choice, selecting a licensing mode in the **Advanced** section that includes Fabric capacity (*Trial*, *Premium*, or *Fabric*).

4. When your new workspace opens, it should be empty.



Create a lakehouse and upload files

Now that you have a workspace, it's time to create a data lakehouse for your data.

1. On the menu bar on the left, select **Create**. In the *New* page, under the *Data Engineering* section, select **Lakehouse**. Give it a unique name of your choice. Make sure the “Lakehouse schemas (Public Preview)” option is disabled.

! **Note:** If the **Create** option is not pinned to the sidebar, you need to select the ellipsis (...) option first.

After a minute or so, a new lakehouse will be created:

2. View the new lakehouse, and note that the **Explorer** pane on the left enables you to browse tables and files in the lakehouse:

You can now ingest data into the lakehouse. There are several ways to do this, but for now you'll download a text file to your local computer (or lab VM if applicable) and then upload it to your lakehouse.

1. Download the [data file](#) from

<https://github.com/MicrosoftLearning/dp-data/raw/main/products.csv>, saving it as *products.csv*.


2. Return to the web browser tab containing your lakehouse, and in the Explorer pane, next to the **Files** folder, select the ... menu. Create a **New subfolder** called *products*.
3. In the ... menu for the *products* folder, **upload** the *products.csv* file from your local computer (or lab VM if applicable).
4. After the file has been uploaded, select the **products** folder to verify that the file has been uploaded, as shown here:

Explore data in a DataFrame

You can now create a Fabric notebook to work with your data. Notebooks provide an interactive environment where you can write and run code.

1. On the menu bar on the left, select **Create**. In the *New* page, under the *Data Engineering* section, select **Notebook**.

A new notebook named **Notebook 1** is created and opened.



2. Fabric assigns a name to each notebook you create, such as Notebook 1, Notebook 2, etc. Click the name panel above the **Home** tab on the menu to change the name to something more descriptive.
3. Select the first cell (which is currently a code cell), and then in the top-right tool bar, use the **M↓** button to convert it to a markdown cell. The text contained in the cell will then be displayed as formatted text.
4. Use the **□ (Edit)** button to switch the cell to editing mode, then modify the markdown as shown below.

Code	Copy
<pre> # Delta Lake tables Use this notebook to explore Delta Lake functionality </pre>	

5. Click anywhere in the notebook outside of the cell to stop editing it.
6. In the **Explorer** pane, select **Add data items**, and then select **Existing data sources**. Connect to the lakehouse you created previously.
7. Add a new code cell, and add the following code to read the products data into a DataFrame using a defined schema:

Code	Copy
<pre> from pyspark.sql.types import StructType, IntegerType, StringType, DoubleType # define the schema schema = StructType() \ .add("ProductID", IntegerType(), True) \ .add("ProductName", StringType(), True) \ .add("Category", StringType(), True) \ .add("ListPrice", DoubleType(), True) df = spark.read.format("csv").option("header","true").schema(schema).load("Files/products/ # df now is a Spark DataFrame containing CSV data from "Files/products/products.csv". display(df) </pre>	

[TIP] Hide or display the explorer panes by using the chevron « icon. This enables you to either focus on the notebook, or your files.

1. Use the **Run cell (>)** button on the left of the cell to run it.

[!NOTE] Since this is the first time you've run any code in this notebook, a Spark session must be started. This means that the first run can take a minute or so to complete. Subsequent runs will be quicker.

1. When the cell code has completed, review the output below the cell, which should look similar to this:

The screenshot shows the Databricks workspace interface. On the left, the 'Explorer' pane displays a 'Data items' section with a 'fab_lakehouse' folder containing 'Tables' and 'Files', and a 'products' file. The main area shows a code cell with the following Python code:

```
from pyspark.sql.types import StructType, IntegerType, StringType, DoubleType
# define the schema
schema = StructType() \
    .add("ProductID", IntegerType(), True) \
    .add("ProductName", StringType(), True) \
    .add("Category", StringType(), True) \
    .add("ListPrice", DoubleType(), True)
df = spark.read.format("csv").option("header", "true").schema(schema).load("Files/products/products.csv")
# df now is a Spark DataFrame containing CSV data from "Files/products/products.csv".
display(df)
```

The code cell output indicates it was executed in 4 seconds. Below the code cell is a 'Table view' section showing the data from the 'products' file:

	ProductID	ProductName	Category	ListPrice
1	771	Mountain-100 Silv...	Mountain Bikes	3399.99
2	772	Mountain-100 Silv...	Mountain Bikes	3399.99
3	773	Mountain-100 Silv...	Mountain Bikes	3399.99
4	774	Mountain-100 Silv...	Mountain Bikes	3399.99
5	775	Mountain-100 Bla...	Mountain Bikes	3374.99
6	776	Mountain-100 Bla...	Mountain Bikes	3374.99
7	777	Mountain-100 Bla...	Mountain Bikes	3374.99

Create Delta tables

You can save the DataFrame as a Delta table by using the `saveAsTable` method. Delta Lake supports the creation of both managed and external tables:

- **Managed** Delta tables benefit from higher performance, as Fabric manages both the schema metadata and the data files.
- **External** tables allow you to store data externally, with the metadata managed by Fabric.

Create a managed table

The data files are created in the **Tables** folder.

1. Under the results returned by the first code cell, use the + Code icon to add a new code cell.

[!TIP] To see the + Code icon, move the mouse to just below and to the left of the output from the current cell.
Alternatively, in the menu bar, on the Edit tab, select + Add code cell.

1. To create a managed Delta table, add a new cell, enter the following code and then run the cell:

Code	
<pre>df.write.format("delta").saveAsTable("managed_products")</pre>	

2. In the Explorer pane, **Refresh** the Tables folder and expand the Tables node to verify that the **managed_products** table has been created.

[!NOTE] The triangle icon next to the file name indicates a Delta table.

The files for managed tables are stored in the **Tables** folder in the lakehouse. A folder named **managed_products** has been created which stores the Parquet files and delta_log folder for the table.

Create an external table

You can also create external tables, which may be stored somewhere other than the lakehouse, with the schema metadata stored in the lakehouse.

1. In the Explorer pane, in the ... menu for the **Files** folder, select **Copy ABFS path**. The ABFS path is the fully qualified path to the lakehouse Files folder.

2. In a new code cell, paste the ABFS path. Add the following code, using cut and paste to insert the abfs_path into the correct place in the code:

Code	 Copy
<pre>df.write.format("delta").saveAsTable("external_products", path="abfs_path/external_products")</pre>	

3. The full path should look similar to this:

Code	 Copy
<pre>abfss://workspace@tenant- onelake.dfs.fabric.microsoft.com/lakehousename.Lakehouse/Files/external_products</pre>	

4. **Run** the cell to save the DataFrame as an external table in the Files/external_products folder.

5. In the Explorer pane, **Refresh** the Tables folder and expand the Tables node and verify that the external_products table has been created containing the schema metadata.

6. In the Explorer pane, in the ... menu for the Files folder, select **Refresh**. Then expand the Files node and verify that the external_products folder has been created for the table's data files.

Compare managed and external tables

Let's explore the differences between managed and external tables using the %%sql magic command.

1. In a new code cell and run the following code:

Code	 Copy
<pre>%%sql DESCRIBE FORMATTED managed_products;</pre>	

2. In the results, view the Location property for the table. Click on the Location value in the Data type column to see the full path. Notice that the OneLake storage location ends with /Tables/managed_products.

3. Modify the DESCRIBE command to show the details of the external_products table as shown here:

Code	 Copy
<pre>%%sql DESCRIBE FORMATTED external_products;</pre>	

4. Run the cell and in the results, view the Location property for the table. Widen the Data type column to see the full path and notice that the OneLake storage locations ends with /Files/external_products.

5. In a new code cell and run the following code:

Code	 Copy
<pre>%%sql DROP TABLE managed_products; DROP TABLE external_products;</pre>	

6. In the Explorer pane, **Refresh** the Tables folder to verify that no tables are listed in the Tables node.
7. In the Explorer pane, **Refresh** the Files folder and verify that the external_products file has *not* been deleted. Select this folder to view the Parquet data files and _delta_log folder.

The metadata for the external table was deleted, but not the data file.

Use SQL to create a Delta table

You will now create a Delta table, using the %%sql magic command.

1. Add another code cell and run the following code:

Code	 Copy
<pre>%%sql CREATE TABLE products USING DELTA LOCATION 'Files/external_products';</pre>	

2. In the Explorer pane, in the ... menu for the **Tables** folder, select **Refresh**. Then expand the Tables node and verify that a new table named *products* is listed. Then expand the table to view the schema.

3. Add another code cell and run the following code:

Code	 Copy
<pre>%%sql SELECT * FROM products;</pre>	

Explore table versioning

Transaction history for Delta tables is stored in JSON files in the delta_log folder. You can use this transaction log to manage data versioning.

1. Add a new code cell to the notebook and run the following code which implements a 10% reduction in the price for mountain bikes:

Code	 Copy
<pre>%%sql UPDATE products SET ListPrice = ListPrice * 0.9 WHERE Category = 'Mountain Bikes';</pre>	

2. Add another code cell and run the following code:

Code	 Copy
<pre>%%sql DESCRIBE HISTORY products;</pre>	

The results show the history of transactions recorded for the table.

1. Add another code cell and run the following code:

Code	 Copy
------	--

```

delta_table_path = 'Files/external_products'
# Get the current data
current_data = spark.read.format("delta").load(delta_table_path)
display(current_data)

# Get the version 0 data
original_data = spark.read.format("delta").option("versionAsOf",
0).load(delta_table_path)
display(original_data)

```

Two result sets are returned - one containing the data after the price reduction, and the other showing the original version of the data.

Analyze Delta table data with SQL queries

Using the SQL magic command you can use SQL syntax instead of Pyspark. Here you will create a temporary view from the products table using a `SELECT` statement.


1. Add a new code cell, and run the following code to create and display the temporary view:

Code	 Copy
<pre> %%sql -- Create a temporary view CREATE OR REPLACE TEMPORARY VIEW products_view AS SELECT Category, COUNT(*) AS NumProducts, MIN(ListPrice) AS MinPrice, MAX(ListPrice) AS MaxPrice, AVG(ListPrice) AS AvgPrice FROM products GROUP BY Category; SELECT * FROM products_view ORDER BY Category; </pre>	

2. Add a new code cell, and run the following code to return the top 10 categories by number of products:

Code	 Copy
<pre> %%sql SELECT Category, NumProducts FROM products_view ORDER BY NumProducts DESC LIMIT 10; </pre>	

3. When the data is returned, select **+ New chart** to display one of the suggested charts.



Alternatively, you can run a SQL query using PySpark.

1. Add a new code cell, and run the following code:

Code Copy

```
from pyspark.sql.functions import col, desc

df_products = spark.sql("SELECT Category, MinPrice, MaxPrice, AvgPrice FROM
products_view").orderBy(col("AvgPrice").desc())
display(df_products.limit(6))
```

Use Delta tables for streaming data

Delta Lake supports streaming data. Delta tables can be a sink or a source for data streams created using the Spark Structured Streaming API. In this example, you'll use a Delta table as a sink for some streaming data in a simulated internet of things (IoT) scenario.

1. Add a new code cell and add the following code and run it:

Code Copy

```

from notebookutils import mssparkutils
from pyspark.sql.types import *
from pyspark.sql.functions import *

# Create a folder
inputPath = 'Files/data/'
mssparkutils.fs.mkdirs(inputPath)

# Create a stream that reads data from the folder, using a JSON schema
jsonSchema = StructType([
    StructField("device", StringType(), False),
    StructField("status", StringType(), False)
])
iotstream = spark.readStream.schema(jsonSchema).option("maxFilesPerTrigger",
1).json(inputPath)

# Write some event data to the folder
device_data = '''{"device":"Dev1","status":"ok"}\n{"device":"Dev1","status":"ok"}\n{"device":"Dev1","status":"ok"}\n{"device":"Dev2","status":"error"}\n{"device":"Dev1","status":"ok"}\n{"device":"Dev1","status":"error"}\n{"device":"Dev2","status":"ok"}\n{"device":"Dev2","status":"error"}\n{"device":"Dev1","status":"ok"}'''

mssparkutils.fs.put(inputPath + "data.txt", device_data, True)

print("Source stream created...")

```

Ensure the message *Source stream created...* is displayed. The code you just ran has created a streaming data source based on a folder to which some data has been saved, representing readings from hypothetical IoT devices.

1. In a new code cell, add and run the following code:

Code	 Copy
<pre> # Write the stream to a delta table delta_stream_table_path = 'Tables/iotdevicedata' checkpointpath = 'Files/delta/checkpoint' deltastream = iotstream.writeStream.format("delta").option("checkpointLocation", checkpointpath).start(delta_stream_table_path) print("Streaming to delta sink...") </pre>	

This code writes the streaming device data in Delta format to a folder named `iotdevicedata`. Because the path for the folder location in the `Tables` folder, a table will automatically be created for it.

1. In a new code cell, add and run the following code:

Code	 Copy
<pre> %%sql SELECT * FROM IotDeviceData; </pre>	

This code queries the `IotDeviceData` table, which contains the device data from the streaming source.

1. In a new code cell, add and run the following code:

```
Code
```

 Copy

```
# Add more data to the source stream
more_data = '''{"device":"Dev1","status":"ok"}
{"device":"Dev1","status":"ok"}
 {"device":"Dev1","status":"ok"}
 {"device":"Dev1","status":"ok"}
 {"device":"Dev1","status":"error"}
 {"device":"Dev2","status":"error"}
 {"device":"Dev1","status":"ok"}'''

mssqlutils.fs.put(inputPath + "more-data.txt", more_data, True)
```

This code writes more hypothetical device data to the streaming source.

1. Re-run the cell containing the following code:

```
Code
```

 Copy

```
%%sql
SELECT * FROM IoTDeviceData;
```

This code queries the IoTDeviceData table again, which should now include the additional data that was added to the streaming source.

1. In a new code cell, add code to stop the stream and run the cell:

```
Code
```

 Copy

```
deltastream.stop()
```

Clean up resources

In this exercise, you've learned how to work with Delta tables in Microsoft Fabric.

If you've finished exploring your lakehouse, you can delete the workspace you created for this exercise.

1. In the bar on the left, select the icon for your workspace to view all of the items it contains.
2. In the ... menu on the toolbar, select **Workspace settings**.
3. In the General section, select **Remove this workspace**.